

CS6510
Applied Machine Learning

Classifiers

12 Aug 2017

Vineeth N Balasubramanian



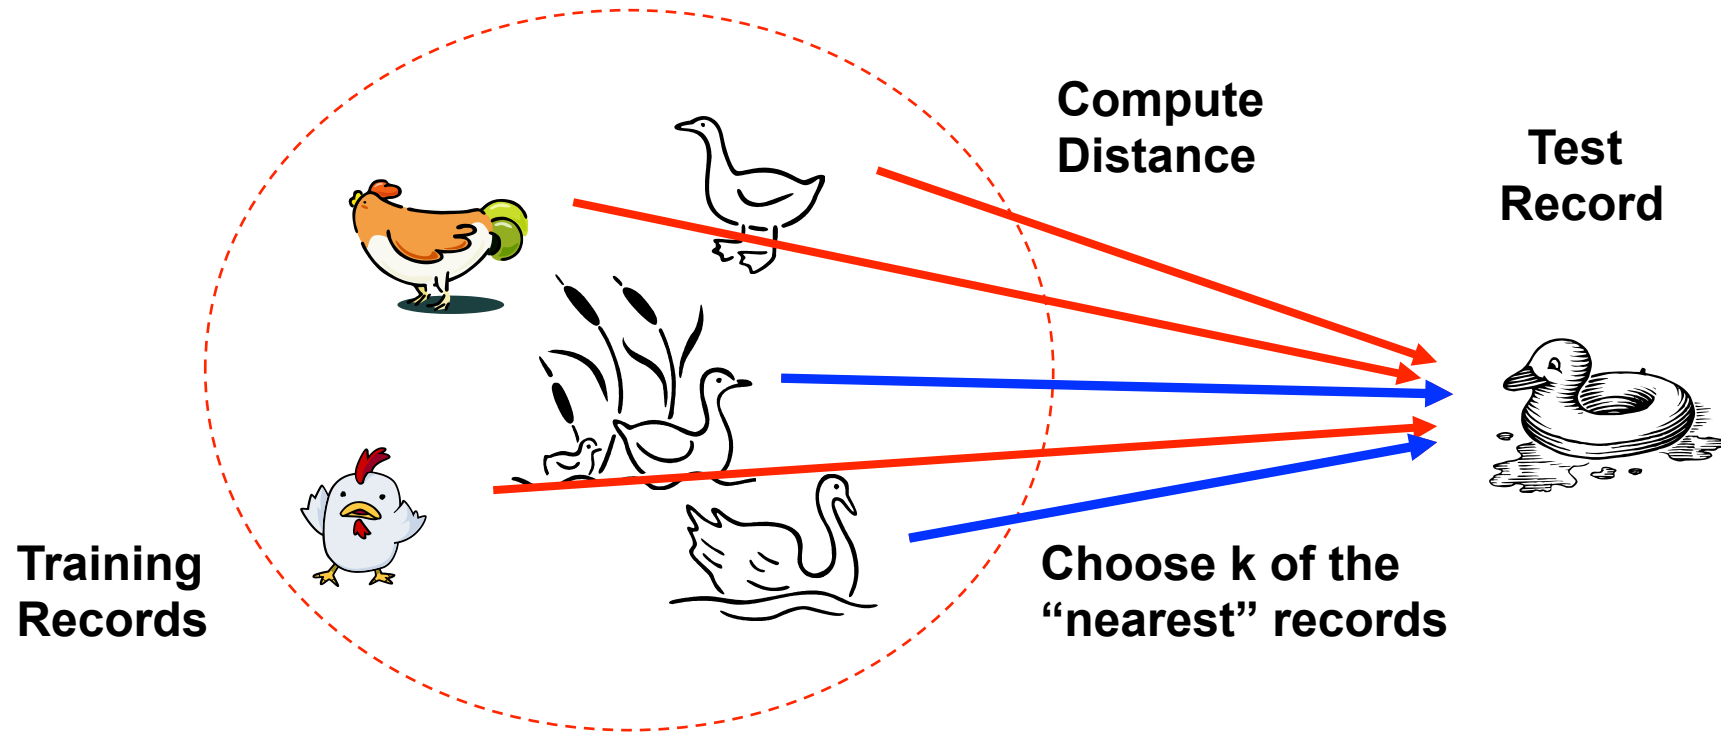
आई आई टी हैदराबाद
IIT Hyderabad

Classification Methods

- **k-Nearest Neighbors**
- Decision Trees
- Naïve Bayes
- Support Vector Machines
- Logistic Regression
- Neural Networks
- Ensemble Methods (Boosting, Random Forests)


k-Nearest Neighbors

- Basic idea:
 - If it walks like a duck, quacks like a duck, then it's probably a duck




k-Nearest Neighbors

- Majority vote within the k nearest neighbors




k-Nearest Neighbors


- Choosing k is important
 - If k is too small, sensitive to noise points
 - If k is too large, neighborhood may include points from other classes



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

k-Nearest Neighbors

- An arbitrary instance is represented by $(a_1(x), a_2(x), a_3(x), \dots, a_n(x))$
 - $a_i(x)$ denotes features
- Euclidean distance between two instances
 - $d(x_i, x_j) = \sqrt{\sum_{r=1}^n (a_r(x_i) - a_r(x_j))^2}$
- In case of continuous-valued target function
 - Mean value of k nearest training examples

How to determine k


- Determined experimentally
 - Start with $k=1$ and use a test set to validate the error rate of the classifier
 - Repeat with $k=k+2$
 - Choose the value of k for which the error rate is minimum
 - Note: k typically an odd number to avoid ties in binary classification

k-Nearest Neighbors

- Eager Learning (**Induction**)
 - Explicit description of target function on the whole training set
- Instance-based Learning (**Transduction**)
 - Learning=storing all training instances
 - Classification=assigning target function to a new instance
 - Referred to as “Lazy” learning

Similar Keywords: K-Nearest Neighbors, Memory-Based Reasoning, Example-Based Reasoning, Instance-Based Learning, Case-Based Reasoning, Lazy Learning

Voronoi Diagram



Decision surface formed by the training examples!

Improvements

- Distance-Weighted Nearest Neighbors
 - Assign weights to the neighbors based on their 'distance' from the query point (E.g., weight 'may' be inverse square of the distances)
- Scaling attributes for fair computation of distances

1 1 1 1 1 1 1 1 1 1 0

0 1 1 1 1 1 1 1 1 1 1

$d = 1.4142$

vs

1 0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0 1

$d = 1.4142$

Solution:
Normalize the
vectors to unit
length

- Measure “closeness” differently
- Finding “close” examples in a large training set quickly
 - E.g. Efficient memory indexing using kd-trees

k-NN: Summary

- Pros

- Highly effective inductive inference method for noisy training data and complex target functions
- Target function for a whole space may be described as a combination of less complex local approximations
- Trains very fast (“Lazy” learner)

- Cons


- Curse of dimensionality
 - In higher dimensions, all the data points lie on the surface of the unit hypersphere (the inside is empty!) Check: <http://www.cs.cmu.edu/~venkatg/teaching/CStheory-infoage/chap1-high-dim-space.pdf>
- Storage: all training examples are saved in memory
 - A decision tree or linear classifier is much smaller
- Slow at query time
 - Can be overcome and presorting and indexing training samples

Convergence of 1-NN

$$P(\text{knnError})$$

$$= 1 - \Pr(y = y_1)$$

$$= 1 - \sum \Pr(Y = y' | x)^2$$



Possible to show that: as the size of training data set approaches infinity, the one nearest neighbor classifier guarantees an error rate of no worse than twice the Bayes error rate (the minimum achievable error rate given the distribution of the data). We will see this later.

k-Nearest Neighbors

- **Parametric** (single global model), **semiparametric** (small number of local models)
- **Nonparametric**: Similar inputs have similar outputs
- Functions (pdf, discriminant, regression) change smoothly
- Keep the training data; “let the data speak for itself”
- Given x , find a small number of closest training instances and interpolate from these

Source: Ethem Alpaydin, *Introduction to Machine Learning*, 3rd Edition (Slides)

Non-parametric Density Estimation

- Given the training set X drawn i.i.d from $p(x)$
- Divide data into bins of size h

- Histogram:
$$\hat{p}(x) = \frac{\# \{x^t \text{ in the same bin as } x\}}{Nh}$$

- Naive estimator:
$$\hat{p}(x) = \frac{\# \{x - h < x^t \leq x + h\}}{2Nh}$$

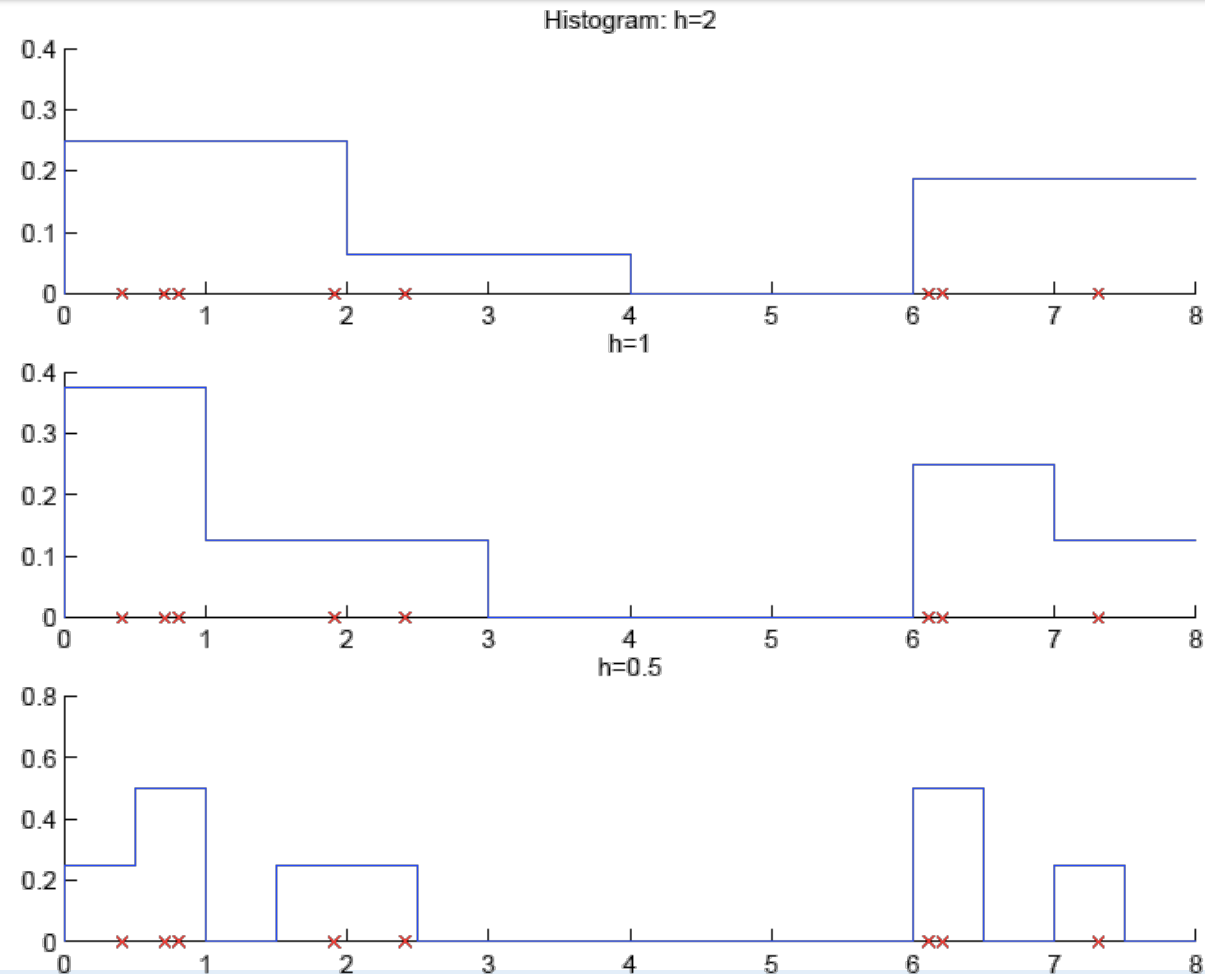
or

$$\hat{p}(x) = \frac{1}{Nh} \sum_{t=1}^N w\left(\frac{x - x^t}{h}\right) \quad w(u) = \begin{cases} 1 & \text{if } |u| < 1/2 \\ 0 & \text{otherwise} \end{cases}$$

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Non-parametric Density Estimation

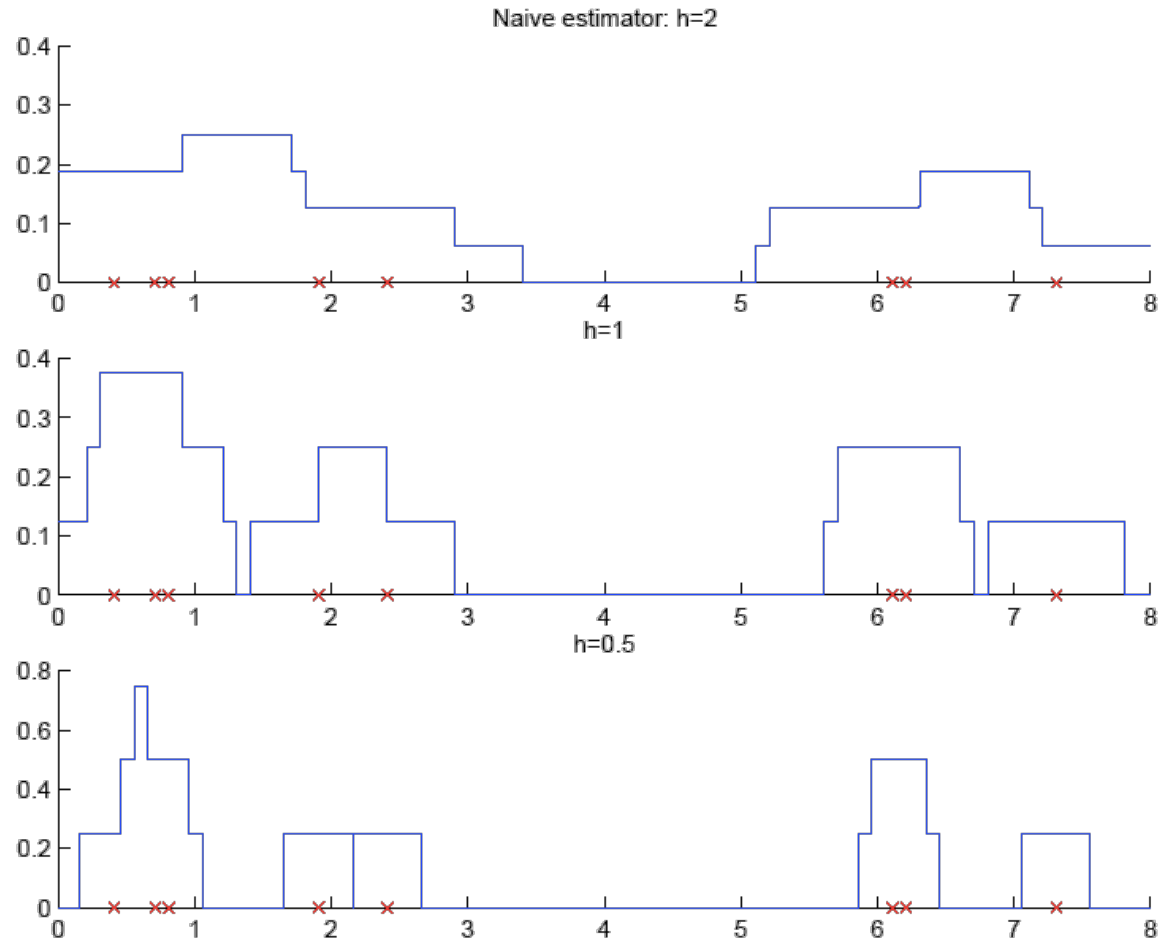
- Histogram Estimation



Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Non-parametric Density Estimation

- Naïve Estimation



Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Non-parametric Density Estimation

- Kernel function, e.g., Gaussian kernel:

$$K(u) = \frac{1}{\sqrt{2\pi}} \exp\left[-\frac{u^2}{2}\right]$$


- Kernel estimator (**Parzen windows**)

$$\hat{p}(x) = \frac{1}{Nh} \sum_{t=1}^N K\left(\frac{x - x^t}{h}\right)$$

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Non-parametric Density Estimation

- Kernel Estimation



Source: Ethem Alpaydin, *Introduction to Machine Learning*, 3rd Edition (Slides)

Non-parametric Density Estimation

- **K-Nearest Neighbor estimator**
- Instead of fixing bin width h and counting the number of instances, fix the instances (neighbors) k and check bin width


$$\hat{p}(x) = \frac{k}{2Nd_k(x)}$$

$d_k(x)$, distance to k th closest instance to x

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Non-parametric Density Estimation

- k-NN Estimation



Source: Ethem Alpaydin, *Introduction to Machine Learning*, 3rd Edition (Slides)

Classification Methods


- k-Nearest Neighbors
- **Decision Trees**
- Naïve Bayes
- Support Vector Machines
- Logistic Regression
- Neural Networks
- Ensemble Methods (Boosting, Random Forests)

Example

PlayTennis: training examples

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Decision Trees



- An efficient nonparametric method
- A hierarchical model
- Divide-and-conquer strategy

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Divide and Conquer

- Internal decision nodes
 - **Univariate:** Uses a single attribute, x_i
 - Numeric x_i :
 - Binary split : $x_i > w_m$
 - Discrete x_i :
 - n -way split for n possible values
 - **Multivariate:** Uses more than one attributes, \mathbf{x}
- Leaves
 - Classification: Class labels, or proportions
 - Regression: Numeric; r average, or local fit
- Learning is **greedy**; find the best split recursively

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)


Classification Trees (C4.5, J48)

- For node m , N_m instances reach m , N_m^i belong to C_i

$$\hat{p}(C_i | \mathbf{x}, m) \equiv p_m^i = \frac{N_m^i}{N_m}$$

- Node m is **pure** if p_m^i is 0 or 1
- Measure of **impurity** is **entropy**

$$I_m = - \sum_{i=1}^K p_m^i \log_2 p_m^i$$



Entropy in information theory specifies the **average (expected) amount of information derived from observing an event**

Source: Ethem Alpaydin, *Introduction to Machine Learning*, 3rd Edition (Slides)

Classification Trees

- If node m is pure, generate a leaf and stop, otherwise split and continue recursively
- **Impurity after split:** N_{mj} of N_m take branch j . N_{mj}^i belong to C_i

$$\hat{P}(C_i | \mathbf{x}, m, j) \equiv p_{mj}^i = \frac{N_{mj}^i}{N_{mj}}$$

$$I'_m = - \sum_{j=1}^n \frac{N_{mj}}{N_m} \sum_{i=1}^K p_{mj}^i \log_2 p_{mj}^i$$

- **Information Gain:** Expected reduction in impurity measure after split
- Choose the *best* attribute(s) (**with maximum information gain**) to split the remaining instances and make that attribute a decision node
 - You can use same logic to find best splitting value too

Source: Ethem Alpaydin, Introduction to Machine Learning, 3rd Edition (Slides)

Other Measures of Impurity


- The properties of functions measuring the **impurity** of a split:
 - $\phi(1/2, 1/2) \geq \phi(p, 1-p)$, for any $p \in [0, 1]$
 - $\phi(0, 1) = \phi(1, 0) = 0$
 - $\phi(p, 1-p)$ is increasing in p on $[0, \frac{1}{2}]$
and decreasing in p on $[\frac{1}{2}, 1]$
- Examples (other than entropy)
 - **Gini impurity/index** $1 - \sum_{j=1}^c p_j^2$

Decision Trees: Example

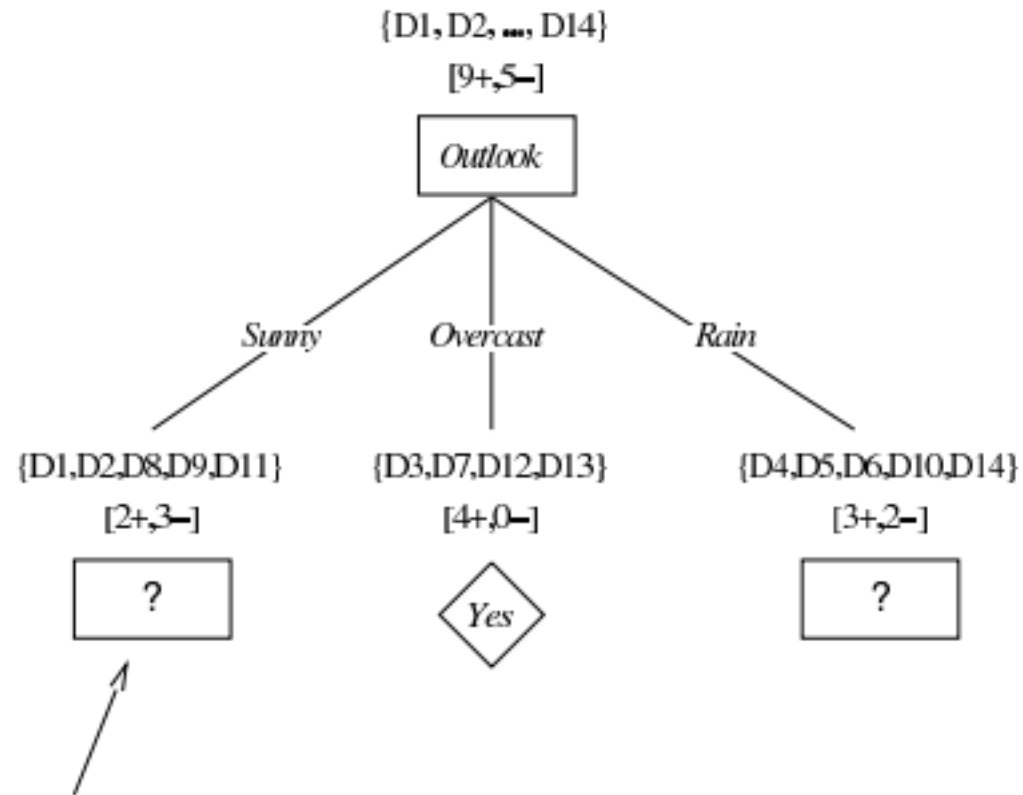
PlayTennis: training examples

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Decision Trees: Example



Decision Trees: Example



Which attribute should be tested here?

$$S_{Sunny} = \{D1, D2, D8, D9, D11\}$$

$$Gain(S_{Sunny}, Humidity) = .970 - (3/5) 0.0 - (2/5) 0.0 = .970$$

$$Gain(S_{Sunny}, Temperature) = .970 - (2/5) 0.0 - (2/5) 1.0 - (1/5) 0.0 = .570$$

$$Gain(S_{Sunny}, Wind) = .970 - (2/5) 1.0 - (3/5) .918 = .019$$

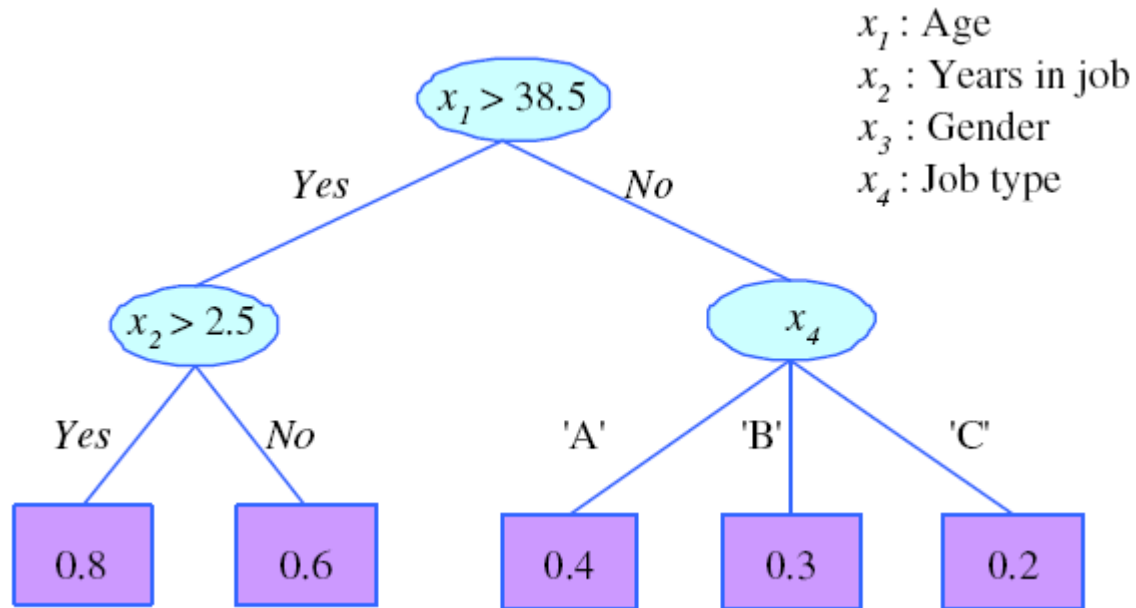
Overfitting and Generalization

- Overfitting can occur with noisy training examples, also when small numbers of examples are associated with leaf nodes.
How to handle?
- **Pruning:** Remove subtrees for better generalization (decrease variance)
 - **Prepruning:** Early stopping
 - **Postpruning:** Grow the whole tree then prune subtrees which overfit on the pruning set
 - Prepruning is faster, postpruning is more accurate

Overfitting and Generalization

- **Occam's Razor principle:** when multiple hypotheses can solve a problem, choose the simplest one
 - a short hypothesis that fits data unlikely to be coincidence
 - a long hypothesis that fits data might be coincidence
- How to select “best” tree:
 - Measure performance over training data
 - Measure performance over separate validation data set
 - **Minimum Description Length:** Minimize $size(tree) + size(misclassifications(tree))$



Rule Extraction from Trees



R1: IF (age>38.5) AND (years-in-job>2.5) THEN $y = 0.8$
R2: IF (age>38.5) AND (years-in-job \leq 2.5) THEN $y = 0.6$
R3: IF (age \leq 38.5) AND (job-type='A') THEN $y = 0.4$
R4: IF (age \leq 38.5) AND (job-type='B') THEN $y = 0.3$
R5: IF (age \leq 38.5) AND (job-type='C') THEN $y = 0.2$

- Convert tree to equivalent set of rules
- Prune each rule independently of others, by removing any preconditions that result in improving its estimated accuracy
- Sort final rules into desired sequence for use

Multivariate Trees



Readings

- Chapters 8, 9, EA Introduction to ML 2nd Edn